

CS 145 Introduction to Databases
Autumn Quarter 2015
Stanford University

Midterm Examination

Tuesday October 27, 2015

80 minutes

	Problem	Full Points	Bonus Points	Your Score
1	The Love of Money	20	10	
2	If It's All the Same to You...	15		
3	ER for the ER	15		
4	Where are my keys?	20		
5	Bob's Bagel Bite Shop	20		
6	Bonus: MVDs		10	
Total		90	20	

The exam contains **21** pages including this cover page. The last page provides you some information that may be useful in the exam. Feel free to tear the last page out.

Instructions:

- This exam is **closed book i.e. no laptops, notes, textbooks, etc. during the exam**
- In all cases, and especially if you're stuck or unsure of your answers, **explain your work!** We'll give partial credit for good explanations of what you were trying to do
- Keep the point values of questions in mind- don't get too stuck on any one part! **Our advice is to skip the bonus problems on your first pass!**

Name: _____

SUNETID: _____

I have read and will abide by the Stanford honor code:

Signature: _____

1 The Love of Money [20 points]

Imagine the following setup:

- Stocks have some **positive or negative** change in price each week
- Traders each pick a portfolio of stocks, i.e. they have **one or more stocks** and they own some **non-zero quantity of each one**

The database schema for this problem consists of two relations corresponding to the descriptions above:

```
Stock(ticker VARCHAR(3), price_change FLOAT, week INT)
Owns(trader_name TEXT, stock_ticker TEXT, quantity FLOAT)
```

Assume there are no NULL values present unless specified otherwise, and that the same stocks are present for all weeks in the database.

1.1 Part 1: Safe for Work [5 points]

Write a query which returns the ticker and price change for all stocks which had a positive price change in week 4:

1.2 Part 2: Portfolio View [5 points]

Recall that our schema is:

```
Stock(ticker VARCHAR(3), price_change FLOAT, week INT)
Owns(trader_name TEXT, stock_ticker TEXT, quantity FLOAT)
```

Now write a query which returns the **ticker** and **net price change** over all the weeks in the table, **for each stock**. Here, for example, if our table covered three weeks, and if stock *ABC* had `price_change=-5.45` in week 1, and `price_change=1.45` in week 2, and `price_change=2.00` in week 3, the *net price change* would be $-5.45 + 1.45 + 2.00 = -2.00$, and we would return a row $(ABC, -2.00)$ for this stock.

1.3 Part 3: Seeking Alpha [10 points]

Recall that our schema is:

```
Stock(ticker VARCHAR(3), price_change FLOAT, week INT)
Owns(trader_name TEXT, stock_ticker TEXT, quantity FLOAT)
```

Write a query which returns the **name of each trader** and **the total change in value of their portfolio** for **week 7**, **only for traders who had a positive total portfolio change**. Here, for example, if the stock *ABC* had `price_change=2.00` in week 7, and the stock *BXZ* had `price_change=-1.50` in week 7, and trader Joe had 3 of *ABC* and 10 of *BXZ*, then the total change in value of Joe's portfolio for week 7 would be $3 * 2 + 10 * (-1.5) = -9.00$. Note that in this example *we would not return a row for Joe since his total portfolio change was ≤ 0* .

1.4 Bonus: Complex Correlated Query [10 points]

Recall that our schema is:

```
Stock(ticker VARCHAR(3), price_change FLOAT, week INT)
Owns(trader_name TEXT, stock_ticker TEXT, quantity FLOAT)
```

No tricks or puns here: this question involves a correlated query! Write a query which returns the **tickers** of stocks that, **in any given week**, out-performed (or matched) all other stocks **which were not in any portfolio with them** that week:

2 If It's All the Same to You... [15 points]

The parts of this problem are **multiple choice**. In each of the below questions, consider the three SQL queries listed, and **circle one of the listed options**. Assume there are no NULL values in any of the tables. The queries will all involve either IN, ALL, ANY (recall, equivalent to "SOME"), or EXISTS.

2.1 Part 1 [5 points]

SELECT S.A FROM S WHERE S.B NOT IN (SELECT R.B FROM R WHERE R.B > 5);

SELECT S.A FROM S WHERE S.B <> ALL (SELECT R.B FROM R WHERE R.B > 5);

SELECT S.A FROM S WHERE NOT EXISTS (SELECT R.B FROM R WHERE R.B > 5 AND R.B = S.B);

- (A) The first query is not equivalent to any of the other queries
- (B) The first query is equivalent to the *second* query only
- (C) The first query is equivalent to the *third* query only
- (D) All three queries are equivalent.

2.2 Part 2 [5 points]

SELECT S.A FROM S WHERE S.B NOT > ANY (SELECT R.B FROM R WHERE R.B > 5);

SELECT S.A FROM S WHERE S.B <= ALL (SELECT R.B FROM R WHERE R.B > 5);

SELECT S.A FROM S WHERE EXISTS (SELECT R.B FROM R WHERE R.B > 5 AND R.B >= S.B);

- (A) The first query is not equivalent to any of the other queries
- (B) The first query is equivalent to the *second* query only
- (C) The first query is equivalent to the *third* query only
- (D) All three queries are equivalent.

2.3 Part 3 [5 points]

SELECT S.A FROM S WHERE S.B IN (SELECT R.B FROM R WHERE R.B > 5);

SELECT S.A FROM S WHERE S.B = ANY (SELECT R.B FROM R WHERE R.B > 5);

SELECT S.A FROM S WHERE S.B >= ALL (SELECT R.B FROM R WHERE R.B > 5);

- (A) The first query is not equivalent to any of the other queries
- (B) The first query is equivalent to the *second* query only
- (C) The first query is equivalent to the *third* query only
- (D) All three queries are equivalent.

3 ER for the ER [15 points]

Design an ER diagram for a hospital that contains the following kinds of objects together with the listed attributes:

- Specialty(name, motto, priority_level)
- Room(room_ID, location, capacity)
- Team(name)
- Doctor(name, SSN, age, sex)

In addition the following must be modeled:

- Each specialty has a unique name (e.g. "cardiology", "neurology", etc.), a motto (for example, rumor has it the Stanford urology group's is "the stream team"), a priority_level, and has **one or more rooms assigned to it**
- **Two specialties cannot use the same room**
- Each specialty may have several **teams**, which can be spread across various rooms (i.e. more than one team from the same specialty can be in the same room). The teams are **defined by their name and their specialty's name**, i.e. "Urology Team 6" is the team with name "Team 6" that belongs to the "Urology" specialty.
- Doctors must **belong to one specialty but may belong to multiple teams within the specialty**
- Doctors may have the same name, but will have different SSNs

Be sure to specify **keys, multiplicity, and participation constraints** in addition to the ER basics (entity sets, relationships and attributes). If in doubt, just explain what you're trying to do!

Extra space:

4 Where are my keys? [20 points]

For the first three parts of this problem, consider the relation $R(U, V, W, X, Y, Z)$ with functional dependencies:

1. $\{W, Y\} \rightarrow \{Z\}$
2. $\{U\} \rightarrow \{X, Y\}$
3. $\{V, X\} \rightarrow \{W\}$
4. $\{Y\} \rightarrow \{V\}$

For all parts of this question: keep in mind that explaining your reasoning / showing your work, while not required, is extremely helpful if we are trying to give you partial credit!

4.1 Part 1 [5 points]

Consider the instance of relation R below. Fill in the blanks so that the instance is consistent with the FDs above. When the FDs imply a specific value for a blank space, fill in that value, otherwise if they do not imply a specific value, fill it in with a question mark:

U	V	W	X	Y	Z
0	0	1	2	3	4
1	0		2	2	
	0	1		3	4
2		1	2	2	5

4.2 Part 2 [5 points]

Find the closure of $\{X, Y\}$, $\{X, Y\}^+$:

4.3 Part 3 [5 points]

List **one** key implied by the given FDs, and **explain why it is a key based on the definition:**

4.4 Part 4 [5 points]

For this problem, consider the relation $T(A, B, C, D, E)$ with the following functional dependencies:

1. $\{A\} \rightarrow \{B, C\}$
2. $\{C\} \rightarrow \{B\}$
3. $\{B\} \rightarrow \{D\}$
4. $\{E\} \rightarrow \{C\}$


Decompose T into two or more tables so that it is in BCNF, and briefly justify that your end result is indeed in BCNF:

5 Bob's Bagel Bite Shop [20 points]

Bob's Bagel Bite Shop, the premiere purveyor of homemade bagel bites in the bay area, has decided to start a loyalty points program, where customers can earn "*bitecoins*". Chef Bob is a big fan of transactions and upholding **ACID** guarantees, but to save on money, he decides to skip buying a full-featured SQL DBMS and instead implement everything himself...

[This space left intentionally blank... draw your favorite animal eating fruit here?]

5.1 Part 1: Write Mostly Ahead? [10 points]



Bob implements the following three procedures, but isn't quite sure how to order them:

1. The values in main memory written by the TXN are flushed to disk
2. The log of the TXN's actions is flushed to disk
3. The commit record (indicating that the TXN was committed) is flushed to disk

5.1.1 [3 points]

Consider the following order: 1,3,2. List the points at which a system crash would cause issues, and explain why, referring to ACID terms (one sentence each):

5.1.2 [3 points]

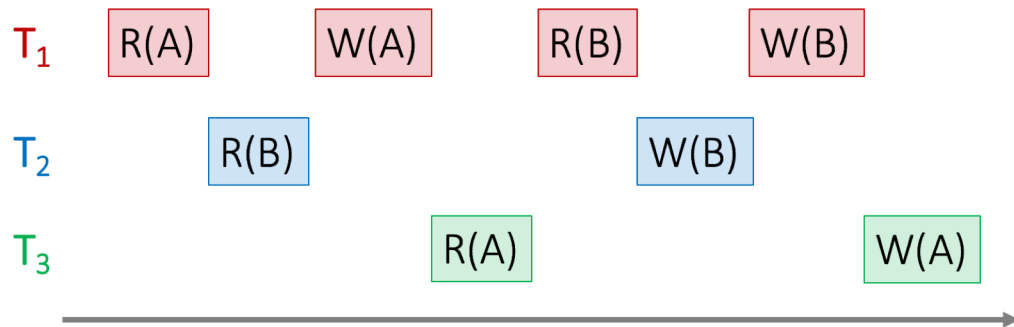
Consider the following order: 2,1,3. List the points at which a system crash would cause issues, and explain why, referring to ACID terms (one sentence each):

5.1.3 [4 points]

What is the best order, and why? Again, refer to ACID terms, and explain with one sentence.

5.2 Part 2: Conflicting Bites [10 points]

Bob also does not have any sort of locking, and must interleave transactions by hand. Consider the time-line below, showing three TXNs operating over two bitcoin accounts, A and B :



5.2.1 [5 points]

What are the conflicts here? List *conflicting* pairs of actions:

5.2.2 [5 points]

Describe the serial order this schedule corresponds to. If this schedule is *not serializable*, first circle two operations whose horizontal position could be swapped to make this schedule serializable.

6 BONUS: MVDs [10 points]

Wow. MVDs. Whaaaaat? Who knew these would be on here. Well, sometimes life throws you for a curve...

Consider the instance of relation R below. Fill in the blanks so that the instance is consistent with the following MVDs:

- $\{U\} \twoheadrightarrow \{W, Y\}$
- $\{V\} \twoheadrightarrow \{X, Z\}$

When the MVDs do not imply a specific value for a blank space, fill it in with a question mark:

U	V	W	X	Y	Z
0		2	3		5
	2		4	5	6
0	2	2	4	4	6
0	1	3	3	5	

[This page left intentionally blank]

7 Possibly Useful Information

- **Canonical SQL Statement:**

```
SELECT <attributes>
FROM <tables>
WHERE <conditions>
GROUP BY <attributes>
HAVING <conditions>
```

- **Functional Dependency (FD):** For a relation R , and sets of attributes X and Y , the functional dependency $X \rightarrow Y$ holds if for any $t_1, t_2 \in R$, $t_1[X] = t_2[X] \implies t_1[Y] = t_2[Y]$.
- **Armstrong's Axioms:** Let the A_i s, B_j s, and C_k s be attributes:
 1. *Split/Combine:* If $A_1, \dots, A_n \rightarrow B_j$ for $j = 1, \dots, m$, then this is equivalent to $A_1, \dots, A_n \rightarrow B_1, \dots, B_m$ and vice-versa
 2. *Reduction/Trivial:* $A_1, \dots, A_n \rightarrow A_i$ for any $i = 1, \dots, n$
 3. *Transitive Closure:* If $A_1, \dots, A_n \rightarrow B_1, \dots, B_m$ and $B_1, \dots, B_m \rightarrow C_1, \dots, C_p$ then $A_1, \dots, A_n \rightarrow C_1, \dots, C_p$
- **Closure:** Given a set of attributes X and a set of FDs F , the closure X^+ is the set of all attributes y such that $X \rightarrow y$.
- **Superkey:** Given a relation R , a superkey is a set of attributes X such that X^+ is equal to the full set of attributes of R .
- **Key:** A key is a minimal superkey, i.e. a superkey where no subset of it is also a superkey.
- **Boyce-Codd Normal Form (BCNF):** A relation R is in BCNF if for all sets of attributes X , either $X^+ = X$ (X is trivial) or $X^+ =$ the set of all attributes (X is a superkey).
- **Conflicts:** Two actions conflict if they are part of different TXNs, involve the same variable, and at least one of them is a write.
- **Serializable:** A schedule is serializable if it is equivalent to some serial ordering.
- **Multi-Value Dependency (MVD):** Given a relation R with a set of attributes A , two sets of attributes $X, Y \subseteq A$, we say that the MVD $X \twoheadrightarrow Y$ holds if for any tuples $t_1, t_2 \in R$ such that $t_1[X] = t_2[X]$, there is a tuple t_3 such that:

- $t_3[X] = t_1[X]$
- $t_3[Y] = t_1[Y]$
- $t_3[A \setminus Y] = t_2[A \setminus Y]$